

Limin Yang

540-998-9158 | limin.yang.cs@gmail.com | <https://whyisyoung.github.io> | GitHub: [whyisyoung](#) | [Google Scholar](#)

SUMMARY

PhD-trained AI security engineer with experience building production security platforms and evaluating autonomous red/blue-team agents. Currently leading agentic-security work at Microsoft, including red-team agent evaluation, security-graph integration, and ground-truth generation from attack logs. Previously led threat-intelligence platform delivery at ByteDance, integrating 150M+ indicators and reducing security-alert false positives at large scale. First-author publications in USENIX Security and IEEE S&P on SOC alerts, ML backdoors, and concept-drift detection.

EXPERIENCE

- Microsoft, Senior Applied AI Engineer, Security AI Team** Dec. 2024 – Current
- Shipped first MVP of an **LLM-based secret-detection** Security Copilot skill under a customer-facing deadline
 - Co-developed an autonomous **red-team multi-agent** system that achieves attack goals without human guidance
 - Built first LLM pipeline turning attack logs into machine checkable ground truth for **red/blue agent evaluations**
 - Built a graph-integration agent that lets custom red-team agents read/write attack activity to an internal security graph, reducing graph ingestion time by 80%
 - Ran custom red-team agents across 30 campaigns; built LLM-based log-evaluation dashboard used by leadership
 - Benchmarked custom red-agent vs. OpenClaw across 11 campaigns; identified gaps justifying guardrail investment
- ByteDance, Senior Software Engineer, Security Operations Center Team** Aug. 2023 – Dec. 2024
- Led the delivery of a multi-source **Threat Intelligence Platform**, integrating 150M+ indicators with a unified schema and aggregation algorithm
 - Saved \$1M+ annually** by evaluating external threat-intel products and designing custom internal threat feeds
 - Protected URLs accessed by **10M+ customers** and delivered actionable intelligence to **10 internal teams**
 - Reduced false positives by 50% in alerts generated by correlating IoCs with HIDS logs at 10+ TB/hour scale
 - Engineered a malware-scanning service that improved detection coverage and streamlined workflows for 3+ teams
- ByteDance, Security Engineering Intern, Threat XDR Team** May 2021 – Aug. 2021
- Enhanced email detection systems, protected **50+ client companies** and detected 5000 additional spams weekly
 - Doubled user-labeled ground-truth dataset through clustering similar emails based on user actions

EDUCATION

- | | |
|--|-----------------------|
| University of Illinois (UIUC), Ph.D. in Computer Science, advisor: Gang Wang | Aug. 2019 – Jul. 2023 |
| Virginia Tech, Ph.D. in Computer Science, advisor: Gang Wang | Aug. 2018 – Jul. 2019 |
| East China Normal University, M.S. Study in Computer Science | Sep. 2015 – Jun. 2018 |
| East China Normal University, B.Eng. in Computer Science | Sep. 2011 – Jun. 2015 |

SELECTED PUBLICATIONS (Citations: 1500+, H-INDEX: 13)

- [USENIX Security'24] **First author**. “An Empirical Measurement of Network Alerts in an SOC”.
- [IEEE S&P'23] **First author**. “Jigsaw Puzzle: Selective Backdoor Attack to Subvert Malware Classifiers”. Code requested by 30+ institutions.
- [USENIX Security'21] **First author**. “CADE: Detecting and Explaining Concept Drift Samples for Security Applications”. **340+ citations, 140+ GitHub stars**.
- [Deep Learning and Security'21] **First author**. “BODMAS: An Open Dataset for Learning based Temporal Analysis of PE Malware”. **270+ citations**. Dataset requested by **140+ institutions**.

TECHNICAL SKILLS

- | | |
|---------------------------|---|
| Languages: | Python, Golang, C++, C, C#, Shell, SQL |
| AI/LLM: | Agents, LLM eval, Prompt Engineering, Agent Harnesses, PyTorch, OpenAI, Claude Code, |
| Cloud & Infra: | Linux, Azure, AWS, Docker, CI/CD, Kubernetes |
| Security: | Agentic Security, Cloud Security, Threat intelligence, Malware detection, Network IDS, Phishing |
| Developer Tools: | Git, Vim, tmux, gRPC, Kafka, MongoDB, Redis, Elasticsearch, Nmap, Metasploit |